


Wishlist saves rarely become purchases — and reviews alone can't explain why

Business goal: increase the % of users who purchase at least one wishlisted item within 30 days of adding it.



Which factor breaks down — and for whom — was unknown going in. That's what this project set out to find, end to end: AI discovery → primary research → a working, evidence-grounded MVP.

An AI discovery engine classified 2,203 user voices across 3 sources — here's how



2,203

reviews classified, checkpointed —
stopped intentionally, not stalled


TEST THE DISCOVERY ENGINE LIVE

[https://zeusworkspace1.app.n8n.cloud/webhook/
wishlist-discovery-engine](https://zeusworkspace1.app.n8n.cloud/webhook/wishlist-discovery-engine)

One honest scope decision: Reddit was excluded. Reddit's Responsible Builder Policy funnels bulk research use through a formal researcher program with a timeline this project's window couldn't accommodate — disclosed here, not quietly worked around.

Reviews found two real barriers — and revealed their own blind spot

Barrier categories in the frozen gold set (n=137)



n = 137

human-labeled gold set — 2 independent labelers, 89.4% raw agreement, 100% category agreement on jointly-flagged items

0.875

classifier accuracy against human labels — measured, not assumed

Four categories never appeared in review corpora at all — not because they're rare, but because reviews are the wrong place to find them. Primary research picks up where the corpus stops.

Every one of 6 interviewees hit the same wall: items disappear before they decide

6/6

interviews independently raised it,
unprompted

32/32

survey respondents (100%) have
experienced it

~0

gold-set corpus examples — reviews
miss it almost entirely

“If something says only two left, I’ll probably decide faster.”

— interview respondent, on how a near-miss stock-out changed future decision speed

Three independent methods — corpus, survey, interviews — converge on one answer: availability decay is real and common, and app reviews are structurally the wrong instrument to measure it.

The real problem: purchase intent, blocked by different uncertainty at the moment of decision

● SEGMENT

Wishlist-savers with unresolved decision-stage uncertainty — spans price, fit, quality, occasion, and availability doubt, not one narrow group.

● ROOT CAUSE

Not disinterest. A specific, often-nameable uncertainty sits between saving and buying — and today's wishlist gives the user no help resolving it.

● OBSERVED WORKAROUND

Users already compensate manually: cross-platform price/review checks, substitute purchases, and rushed decisions after a stock-out scare.

● WHY IT MATTERS

User: less regret, less re-research effort. Business: converts demand that's already explicit and already on the platform — no acquisition cost.

A Discovery Concierge that resolves the one barrier blocking each save

Constraint honored: no monetary incentives — every nudge resolves real uncertainty. None manufacture urgency or offer a discount.

1

Evidence-backed reasoning

Every recommendation traces to a real theme, a real review, or a real interview finding — never a generic “picked for you.”

2

Honest confidence states


High / Medium / Insufficient-evidence, always icon + color + label together — confidence is shown, never implied by color alone.

3

No manufactured urgency

“What if I wait” tells the truth — a real scarcity fact, or an honest “nothing changes.” Never a fake countdown.

Live MVP: real evidence, never guessed



FLAGSHIP CASE — WIDE-LEG JEANS

“If it was 1,200 I probably would've bought it then and there.” — real interview quote, directly powering the item's price-history chart and confidence reasoning in the live MVP. Not a template filled with a plausible-sounding guess.

LIVE MVP

<https://mvp-henna-delta.vercel.app>

We don't trust AI insights blindly

0.875

E1 — classifier accuracy measured against 137 human-labeled reviews, not self-graded.

Kappa = 0.159

E3 — cross-model agreement, disclosed as a real limitation of the smaller local backup model, not smoothed over.

The categories aren't the problem — the small local model is.

Two independent human labelers, using the same taxonomy, reached 89.4% raw agreement and 100% category agreement on items they both flagged relevant. The task is learnable. E3's low kappa reflects a genuinely undersized local model standing in for a stronger one — not an ambiguous codebook. Production classification uses the stronger hosted model; the small local model is a documented fallback, not the primary path.

Success means saves that convert, not saves that pile up

● NORTH STAR

Wishlist-to-Purchase, Within 30 Days

The business metric this whole project decomposes and targets. Unchanged.

● MVP LEADING INDICATORS

Trace-widget open rate · Trace-to-cart conversion

Does the reasoning get read — and does resolving uncertainty actually lead to action?

● GUARDRAIL

False-urgency rate

“What if I wait” shown where stock/price didn't actually change afterward. Designed in from day one, not added for this slide.

What could go wrong, and what we're doing about it

Thin categories, wobbly numbers

fit_size / quality_trust / price_certainty rest on small samples (n=5-11); recall moved run to run (0.4→0.5→0.545).

MITIGATION: Treated as directional, disclosed in every writeup — never presented as exact.

Corpus skews toward complaints

Review-writers aren't representative; Reddit was scoped out for policy reasons.

MITIGATION: Triangulated against survey (n=32) and interviews (n=6) — caught what the corpus alone would have missed.

Local-model consistency is weak

$\kappa = 0.159$ between the primary and backup classifiers.

MITIGATION: Production path uses the stronger hosted model with retry-escalation; local is a documented fallback only.

Demo reasoning is templated

The MVP doesn't yet call a live model against real per-user history.

MITIGATION: Disclosed openly. Deterministic-retrieval-then-narration is already the intended production architecture.