



Wishlist saves rarely become purchases — and reviews alone can't explain why

Business goal: increase the % of users who purchase at least one wishlisted item within 30 days of adding it.



Which factor breaks down — and for whom — was unknown going in. That's what this project set out to find, end to end: AI discovery → primary research → a working, evidence-grounded MVP.

An AI discovery engine classified 2,203 user voices across 3 sources — here's how



2,203


reviews classified, checkpointed —
stopped intentionally, not stalled

TEST THE DISCOVERY ENGINE LIVE

[https://zeusworkspace1.app.n8n.cloud/webhook/
wishlist-discovery-engine](https://zeusworkspace1.app.n8n.cloud/webhook/wishlist-discovery-engine)

Sources actually scraped at scale: Play Store + App Store reviews. Two were consciously not pursued in this window, not silently skipped: Reddit was scoped out per Reddit's Responsible Builder Policy, which funnels bulk research use through a formal researcher program this project's timeline couldn't accommodate. YouTube comments and social/community sources were never attempted — disclosed here for the same reason, not hidden.

Reviews found two real barriers — and revealed their own blind spot



n = 137

human-labeled gold set — 2 independent labelers,
89.4% raw agreement, 100% category agreement
on jointly-flagged items

0.875

classifier accuracy against human labels —
measured, not assumed

Four categories never appeared in review corpora at all — not because they're rare, but because reviews are the wrong place to find them. Primary research picks up where the corpus stops.

Every one of 6 interviewees hit the same wall: items disappear before they decide

6/6

interviews independently raised it,
unprompted

32/32

survey respondents (100%) have
experienced it

~0

gold-set corpus examples — reviews
miss it almost entirely

“If something says only two left, I’ll probably decide faster.”

— interview respondent, on how a near-miss stock-out changed future decision speed

Three independent methods — corpus, survey, interviews — converge on one answer: availability decay is real and common, and app reviews are structurally the wrong instrument to measure it.

SURVEY (n=32) docs.google.com/forms/d/e/1FAIpQLScsU5OcvTbEetugF0pG-ek2eo_8VEz1DLEjX0gMDsydgOHizA/viewform

The real problem: purchase intent, blocked by different uncertainty at the moment of decision

SEGMENT

1

Wishlist-savers with unresolved decision-stage uncertainty — spans price, fit, quality, occasion, and availability doubt, not one narrow group.

ROOT CAUSE

2

Not disinterest. A specific, often-nameable uncertainty sits between saving and buying — today's wishlist gives no help resolving it.

OBSERVED WORKAROUND

3

Users already compensate manually: cross-platform price/review checks, substitute purchases, rushed decisions after a stock-out scare.

WHY IT MATTERS

4

User: less regret, less re-research effort. Business: converts demand that's already explicit and already on the platform — no acquisition cost.

Process note: the brief's kill-criteria single-segment selection could not be run — no segment-tagged corpus data exists, and building it risked reopening the frozen gold set or inferring segments never observed. Full reasoning: docs/decisions/opportunity-selection.md

A Discovery Concierge that resolves the one barrier blocking each save

Constraint honored: no monetary incentives — every nudge resolves real uncertainty. None manufacture urgency or offer a discount.

1

Evidence-backed reasoning

Every recommendation traces to a real theme, a real review, or a real interview finding — never a generic “picked for you.”

2

Honest confidence states

High / Medium / Insufficient-evidence, always icon + color + label together — confidence is shown, never implied by color alone.


3

No manufactured urgency

“What if I wait” tells the truth — a real scarcity fact, or an honest “nothing changes.” Never a fake countdown.

Also reachable — Compare Similar (a real side-by-side against another saved item, never shown empty) and Alternatives (secondary browse, substitutes tied to existing wishlist items). Same 3 rules apply: honest CTAs only — Save / Compare / Move to cart, no “Buy Now” anywhere.

Live MVP: real evidence, never guessed




FLAGSHIP CASE — WIDE-LEG JEANS

"If it was 1,200 I probably would've bought it then and there." — real interview quote, directly powering the item's price-history box and confidence reasoning above. Not a template filled with a plausible-sounding guess.

LIVE MVP <https://mvp-henna-delta.vercel.app>

We don't trust AI insights blindly



The categories aren't the problem — the small local model is.

Two independent human labelers, using the same taxonomy, reached 89.4% raw agreement and 100% category agreement on items they both flagged relevant. The task is learnable. E3's low kappa reflects a genuinely undersized local model standing in for a stronger one — not an ambiguous codebook. Production classification uses the stronger hosted model; the small local model is a documented fallback, not the primary path.

FULL REPO & RESEARCH TRAIL github.com/blazing-llama/NL-Myntra-Graduation-Project

Success means saves that convert, not saves that pile up

NORTH STAR

Wishlist-to-Purchase, Within 30 Days

The business metric this whole project decomposes and targets. Unchanged.

MVP LEADING INDICATORS

Trace-widget open rate · Trace-to-cart conversion

Does the reasoning get read — and does resolving uncertainty actually lead to action?

GUARDRAIL

False-urgency rate

“What if I wait” shown where stock/price didn't actually change afterward. Designed in from day one, not added for this slide.

What could go wrong, and what we're doing about it

Thin categories, wobbly numbers

1

fit_size / quality_trust / price_certainty rest on small samples (n=5–11); recall moved run to run (0.4→0.5→0.545).

MITIGATION: Treated as directional, disclosed in every writeup — never presented as exact.

Corpus skews toward complaints

2

Review-writers aren't representative; Reddit/YouTube/social scoped out for policy/time reasons.

MITIGATION: Triangulated against survey (n=32) and interviews (n=6) — caught what the corpus alone would have missed.

Local-model consistency is weak

3

Kappa = 0.159 between the primary and backup classifiers.

MITIGATION: Production path uses the stronger hosted model with retry-escalation; local is a documented fallback only.

Demo reasoning is templated

4

The MVP doesn't yet call a live model against real per-user history.

MITIGATION: Disclosed openly. Deterministic-retrieval-then-narration is already the intended production architecture.